

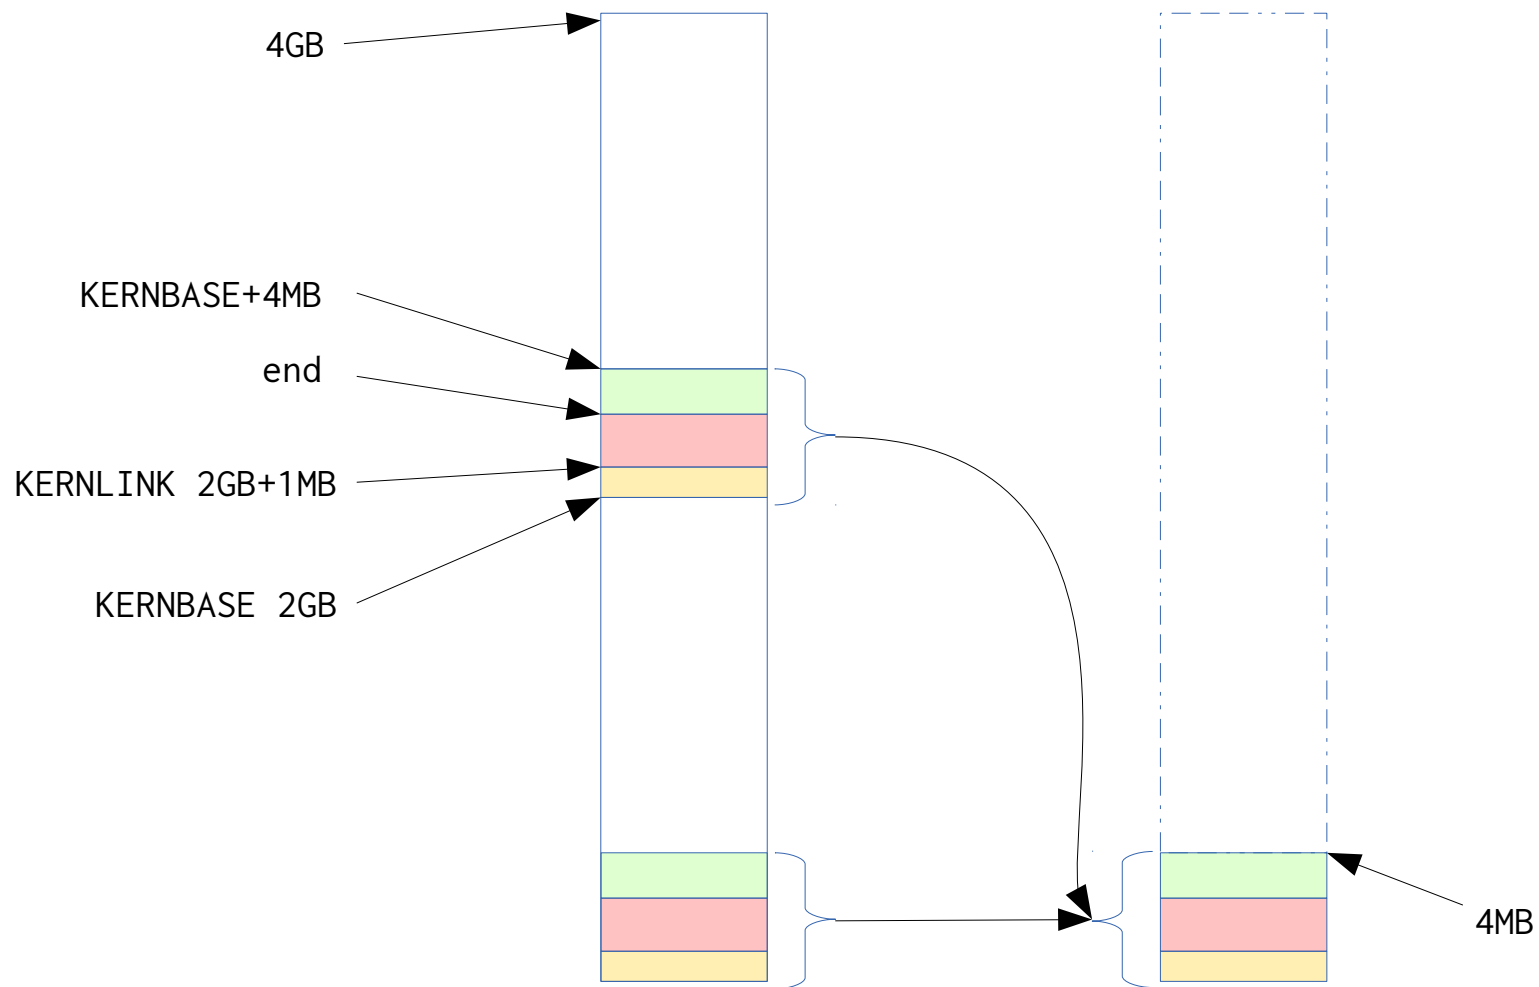
Operativni sistemi

dr.sc. Amer Hasanović

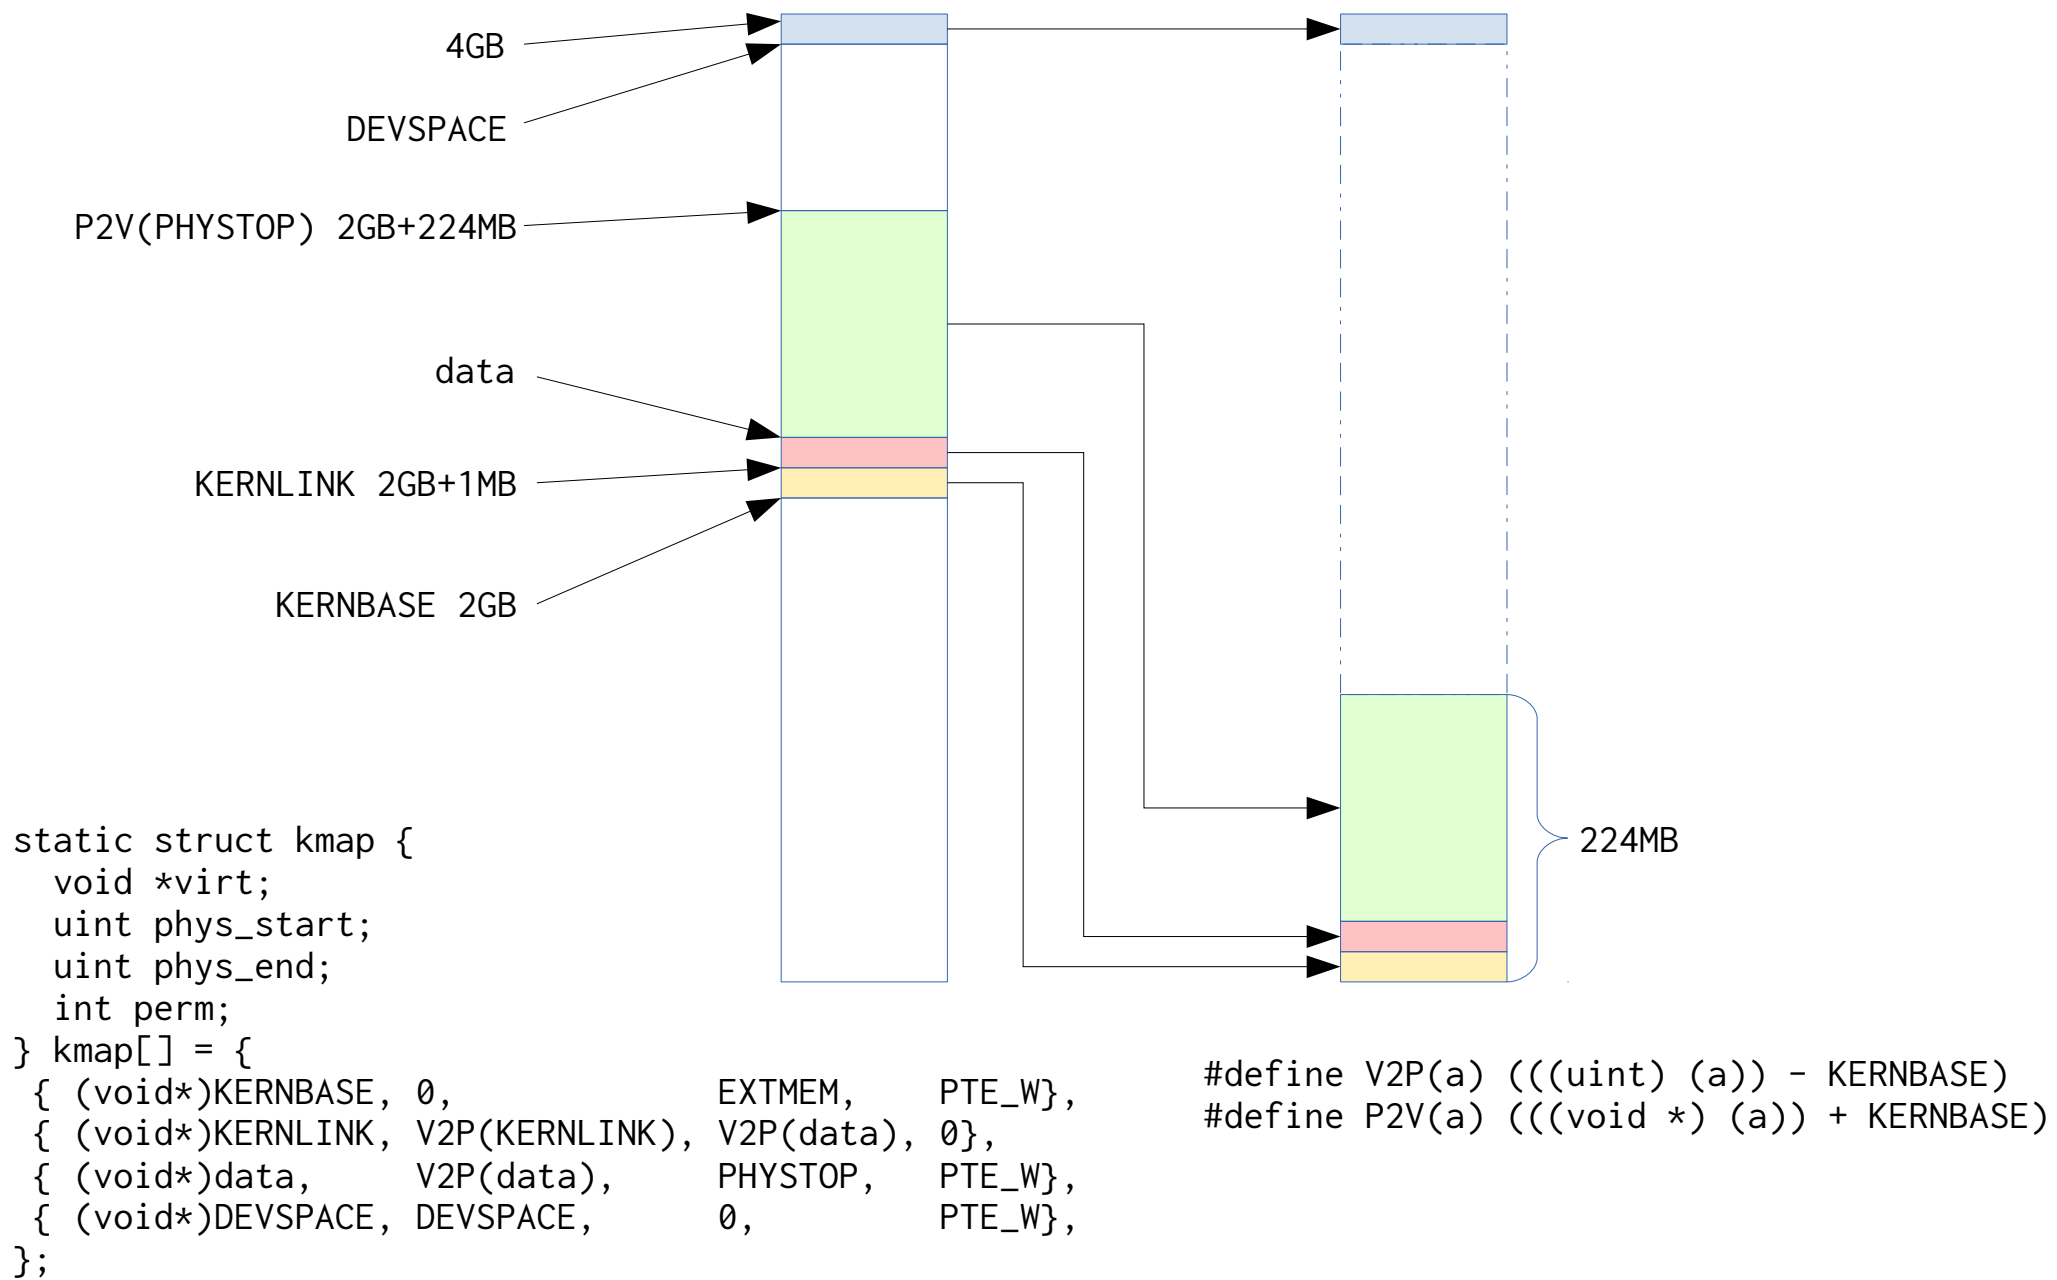
Inicijalni kernel virtuelni adresni (VA) prostor

- Nakon preuzimanja kontrole od bootloader-a, OS formira virtuelni adresni prostor na način da aktivira jednostavno straničenje:
 - dvije 4MB stranice iz virtuelne memorije mapiraju se na istu fizičku lokaciju:
 - $[0, 4\text{MB})$ VA $\rightarrow [0, 4\text{MB})$ PA
 - $[2\text{GB}, 2\text{GB}+4\text{MB})$ VA $\rightarrow [0, 4\text{MB})$ PA
 - niz `entrypgdir` sadrži detalje ovog mapiranja.
- Formirano mapiranje koristi se samo tokom boot procesa.
- Da bi mogao da efikasno podrži više paralelnih procesa kao i da alocira prostor za potrebne data strukture, kernel mora kreirati novi virtuelni adresni prostor baziran na stranicama veličine 4KB.

Inicijalno VA → PA mapiranje



Kernel adresni prostor



- Da bi se kreirao željeni virtuelni adresni prostor potrebno je:
 - podijeliti 4 segmenta VA prostora na stranice veličine po 4KB
 - formirati MMU mapiranje VA → PA straničenjem:
 - kreirati i adevatno konfigurirati tabelu direktorija tj PD:
 - PD staje u jednu stranicu od 4KB
 - kreirati i adekvatno konfigurirati tabele za translaciju stranica (PT) u aktivnim direktorijima:
 - i to, po jedan PT za svaki aktivni direktorij;
 - jedan PT staje u jednu stranicu od 4KB.
 - učitati PD adresu u %cr3 nakon čega CPU počinje da koristi novokreirani VA prostor

Alokator stranica

- Data struktura koja upravlja prostorom koji je neophodan za pohranu PD i PT-ova, kao i ostalih objekata koje dinamički koristi kernel:
 - inicijalno koristi slobodni prostor u intervalu $[end, KERNBASE+4MB)$, preostao nakon učitavanja kernela u memoriju (funkcija `kinit1` aktivira ovaj slobodni prostor)
 - pred kraj boot procesa upravljani prostor proširuje se do lokacije $KERNBASE+PHYSTOP$; (funkcija `kinit2` aktivira ovaj slobodni prostor)
 - organizira slobodan prostor kojim upravlja u manje komade memorije od po 4KB (stranice).

- održava informaciju o slobodnim stranicama pomoću posebne linkane liste slobodnih stranica:
 - pointer na početak aktuelne liste slobodnih stranica drži se u globalnoj strukturi `kmem`;
 - na početku svake slobodne stranice nalazi se pointer na početak slijedeće slobodne stranice;
 - funkcije `kalloc` i `kfree` manipuliraju ovom strukturom.


```


struct {
    struct spinlock lock;
    int use_lock;
    struct run *freelist;
} kmem;

```

```

struct run {
    struct run *next;
};

```



```

void kfree(char *v)
{
    struct run *r;
    memset(v, 1, PGSIZE);
    r = (struct run*)v;
    r->next = kmem.freelist;
    kmem.freelist = r;
}

```

```

char* kalloc(void)
{
    struct run *r;
    r = kmem.freelist;
    if(r)
        kmem.freelist = r->next;
    return (char*)r;
}

```

```


struct {
    struct spinlock lock;
    int use_lock;
    struct run *freelist;
} kmem;

```

```

struct run {
    struct run *next;
};

```



```

void kfree(char *v)
{
    struct run *r;
    memset(v, 1, PGSIZE);
    r = (struct run*)v;
    r->next = kmem.freelist;
    kmem.freelist = r;
}

```

Prvi poziv: kfree(end);

```


struct {
    struct spinlock lock;
    int use_lock;
    struct run *freelist;
} kmem;

```

```

struct run {
    struct run *next;
};

```



```

void kfree(char *v)
{
    struct run *r;
    memset(v, 1, PGSIZE);
    r = (struct run*)v;
    r->next = kmem.freelist;
    kmem.freelist = r;
}

```

Drugi poziv: `kfree(end+PGSIZE);`

```


struct {
    struct spinlock lock;
    int use_lock;
    struct run *freelist;
} kmem;


```

```

struct run {
    struct run *next;
};

```





```
kinit1(end, P2V(4*1024*1024));
```

```
void kinit1(void *vstart, void *vend)
{
    //...
    freerange(vstart, vend);
}
```

```
void freerange(void *vstart, void *vend)
{
    char *p;
    p = (char*)PGROUNDUP((uint)vstart);
    for(; p + PGSIZE <= (char*)vend; p += PGSIZE)
        kfree(p);
}
```

Formiranje kernel AP

```

int main(void)
{
    kinit1(end, P2V(4*1024*1024)); // phys page allocator
    kvmalloc(); // kernel page table
    mpinit(); // collect info about this machine
    lapicinit();
    seginit(); // set up segments
    cprintf("\ncpu%d: starting xv6\n\n", cpu->id);
    picinit(); // interrupt controller
    ioapicinit(); // another interrupt controller
    consoleinit(); // I/O devices & their interrupts
    uartinit(); // serial port
    pinit(); // process table
    tvinit(); // trap vectors
    binit(); // buffer cache
    fileinit(); // file table
    iinit(); // inode cache
    ideinit(); // disk
    if(!ismp)
        timerinit(); // uniprocessor timer
    startothers(); // start other processors
    kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
    userinit(); // first user process
    // Finish setting up this processor in mpmain.
    mpmain();
}

```

```

int main(void)
{
    kinit1(end, P2V(4*1024*1024)); // phys page allocator
    kvmalloc(); // kernel page table
    mpinit(); // collect info about this machine
    lapicinit();
    seginit(); // set up segments
    cprintf("\ncpu%d: starting xv6\n\n", cpu->id);
    picinit(); // interrupt controller
    ioapicinit(); // another interrupt controller
    consoleinit(); // I/O devices & their interrupts
    uartinit(); // serial port
    pinit(); // process table
    tvinit(); // trap vectors
    binit(); // buffer cache
    fileinit(); // file table
    iinit(); // inode cache
    ideinit(); // disk
    if(!ismp)
        timerinit(); // uniprocessor timer
    startothers(); // start other processors
    kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
    userinit(); // first user process
    // Finish setting up this processor in mpmain.
    mpmain();
}

```



```
void kvmalloc(void)
{
    kpgdir = setupkvm();
    switchkvm();
}
```

```
pde_t* setupkvm(void)
```

```
{
    pde_t *pgdir;
    struct kmap *k;

    if((pgdir = (pde_t*)kalloc()) == 0)
        return 0;

    memset(pgdir, 0, PGSIZE);
    for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
        if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
                    (uint)k->phys_start, k->perm) < 0)
            return 0;
    return pgdir;
}
```



```
void kvmalloc(void)
```

```
{
    kpgdir = setupkvm();
    switchkvm();
}
```

```
static struct kmap {
```

```
    void *virt;
```

```
    uint phys_start;
```

```
    uint phys_end;
```

```
    int perm;
```

```
} kmap[] = {
```

```
    { (void*)KERNBASE, 0,          EXTMEM,      PTE_W},
```

```
    { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0},
```

```
    { (void*)data,      V2P(data),    PHYSTOP,    PTE_W},
```

```
    { (void*)DEVSPACE, DEVSPACE,      0,          PTE_W},
```

```
};
```

```

pde_t* setupkvm(void)
{
    pde_t *pgdir;
    struct kmap *k;

    if((pgdir = (pde_t*)kalloc()) == 0)
        return 0;

    memset(pgdir, 0, PGSIZE);
    for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
        if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
                    (uint)k->phys_start, k->perm) < 0)
            return 0;
    return pgdir;
}


```



```

void kvmalloc(void)
{
    kpgdir = setupkvm();
    switchkvm();
}

```



```

static int mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
{
    char *a, *last;
    pte_t *pte;

    a = (char*)PGROUNDDOWN((uint)va);
    last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
    for(;;){
        if((pte = walkpgdir(pgdir, a, 1)) == 0)
            return -1;
        if(*pte & PTE_P)
            panic("remap");
        *pte = pa | perm | PTE_P;
        if(a == last)
            break;
        a += PGSIZE;
        pa += PGSIZE;
    }
    return 0;
}

```

```

static pte_t * walkpgdir(pde_t *pgdir, const void *va, int alloc)
{
    pde_t *pde;
    pte_t *pgtab;

    pde = &pgdir[PDX(va)];

    if(*pde & PTE_P)
    {
        pgtab = (pte_t*)p2v(PTE_ADDR(*pde));
    }
    else
    {
        if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
            return 0;

        // Make sure all those PTE_P bits are zero.

        memset(pgtab, 0, PGSIZE);

        // The permissions here are overly generous, but they can
        // be further restricted by the permissions in the page table
        // entries, if necessary.

        *pde = v2p(pgtab) | PTE_P | PTE_W | PTE_U;
    }
    return &pgtab[PTX(va)];
}

```

```

pde_t* setupkvm(void)
{
    pde_t *pgdir;
    struct kmap *k;

    if((pgdir = (pde_t*)kalloc()) == 0)
        return 0;

    memset(pgdir, 0, PGSIZE);
    for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
        if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
                    (uint)k->phys_start, k->perm) < 0)
            return 0;
    return pgdir;
}


```



```

void kvmalloc(void)
{
    kpgdir = setupkvm();
    switchkvm();
}


```



```

void switchkvm(void)
{
    lcr3(v2p(kpgdir));
}

```



```

static int mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
{
    char *a, *last;
    pte_t *pte;

    a = (char*)PGROUNDDOWN((uint)va);
    last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
    for(;;){
        if((pte = walkpgdir(pgdir, a, 1)) == 0)
            return -1;
        if(*pte & PTE_P)
            panic("remap");
        *pte = pa | perm | PTE_P;
        if(a == last)
            break;
        a += PGSIZE;
        pa += PGSIZE;
    }
    return 0;
}

```