

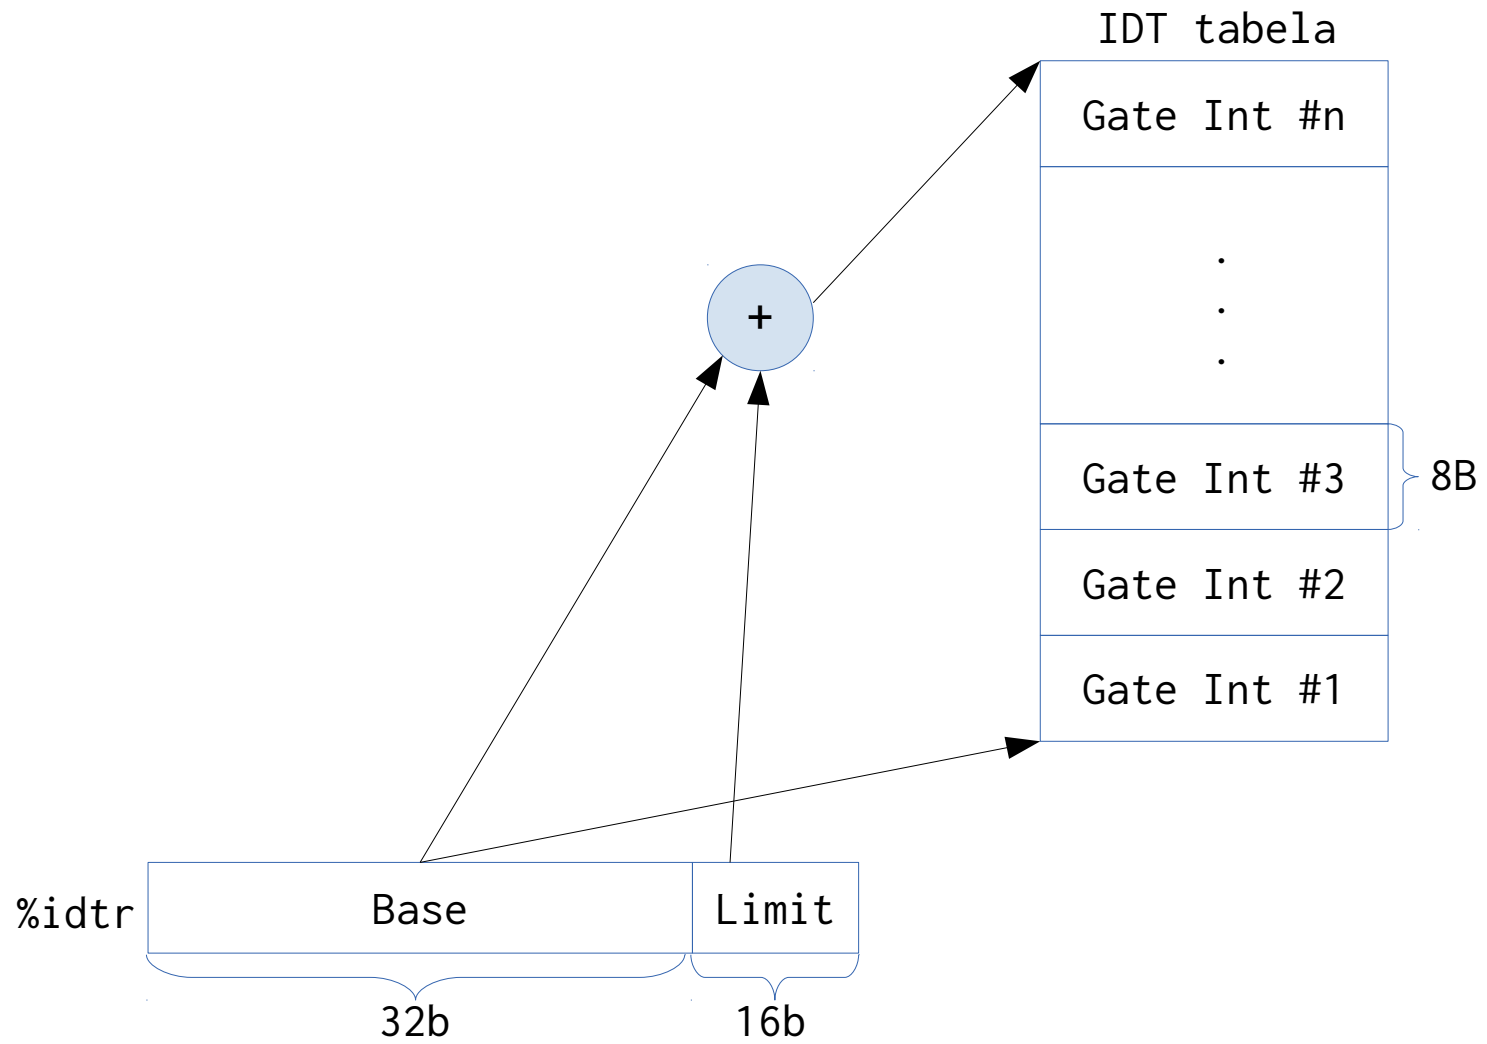
# Operativni sistemi

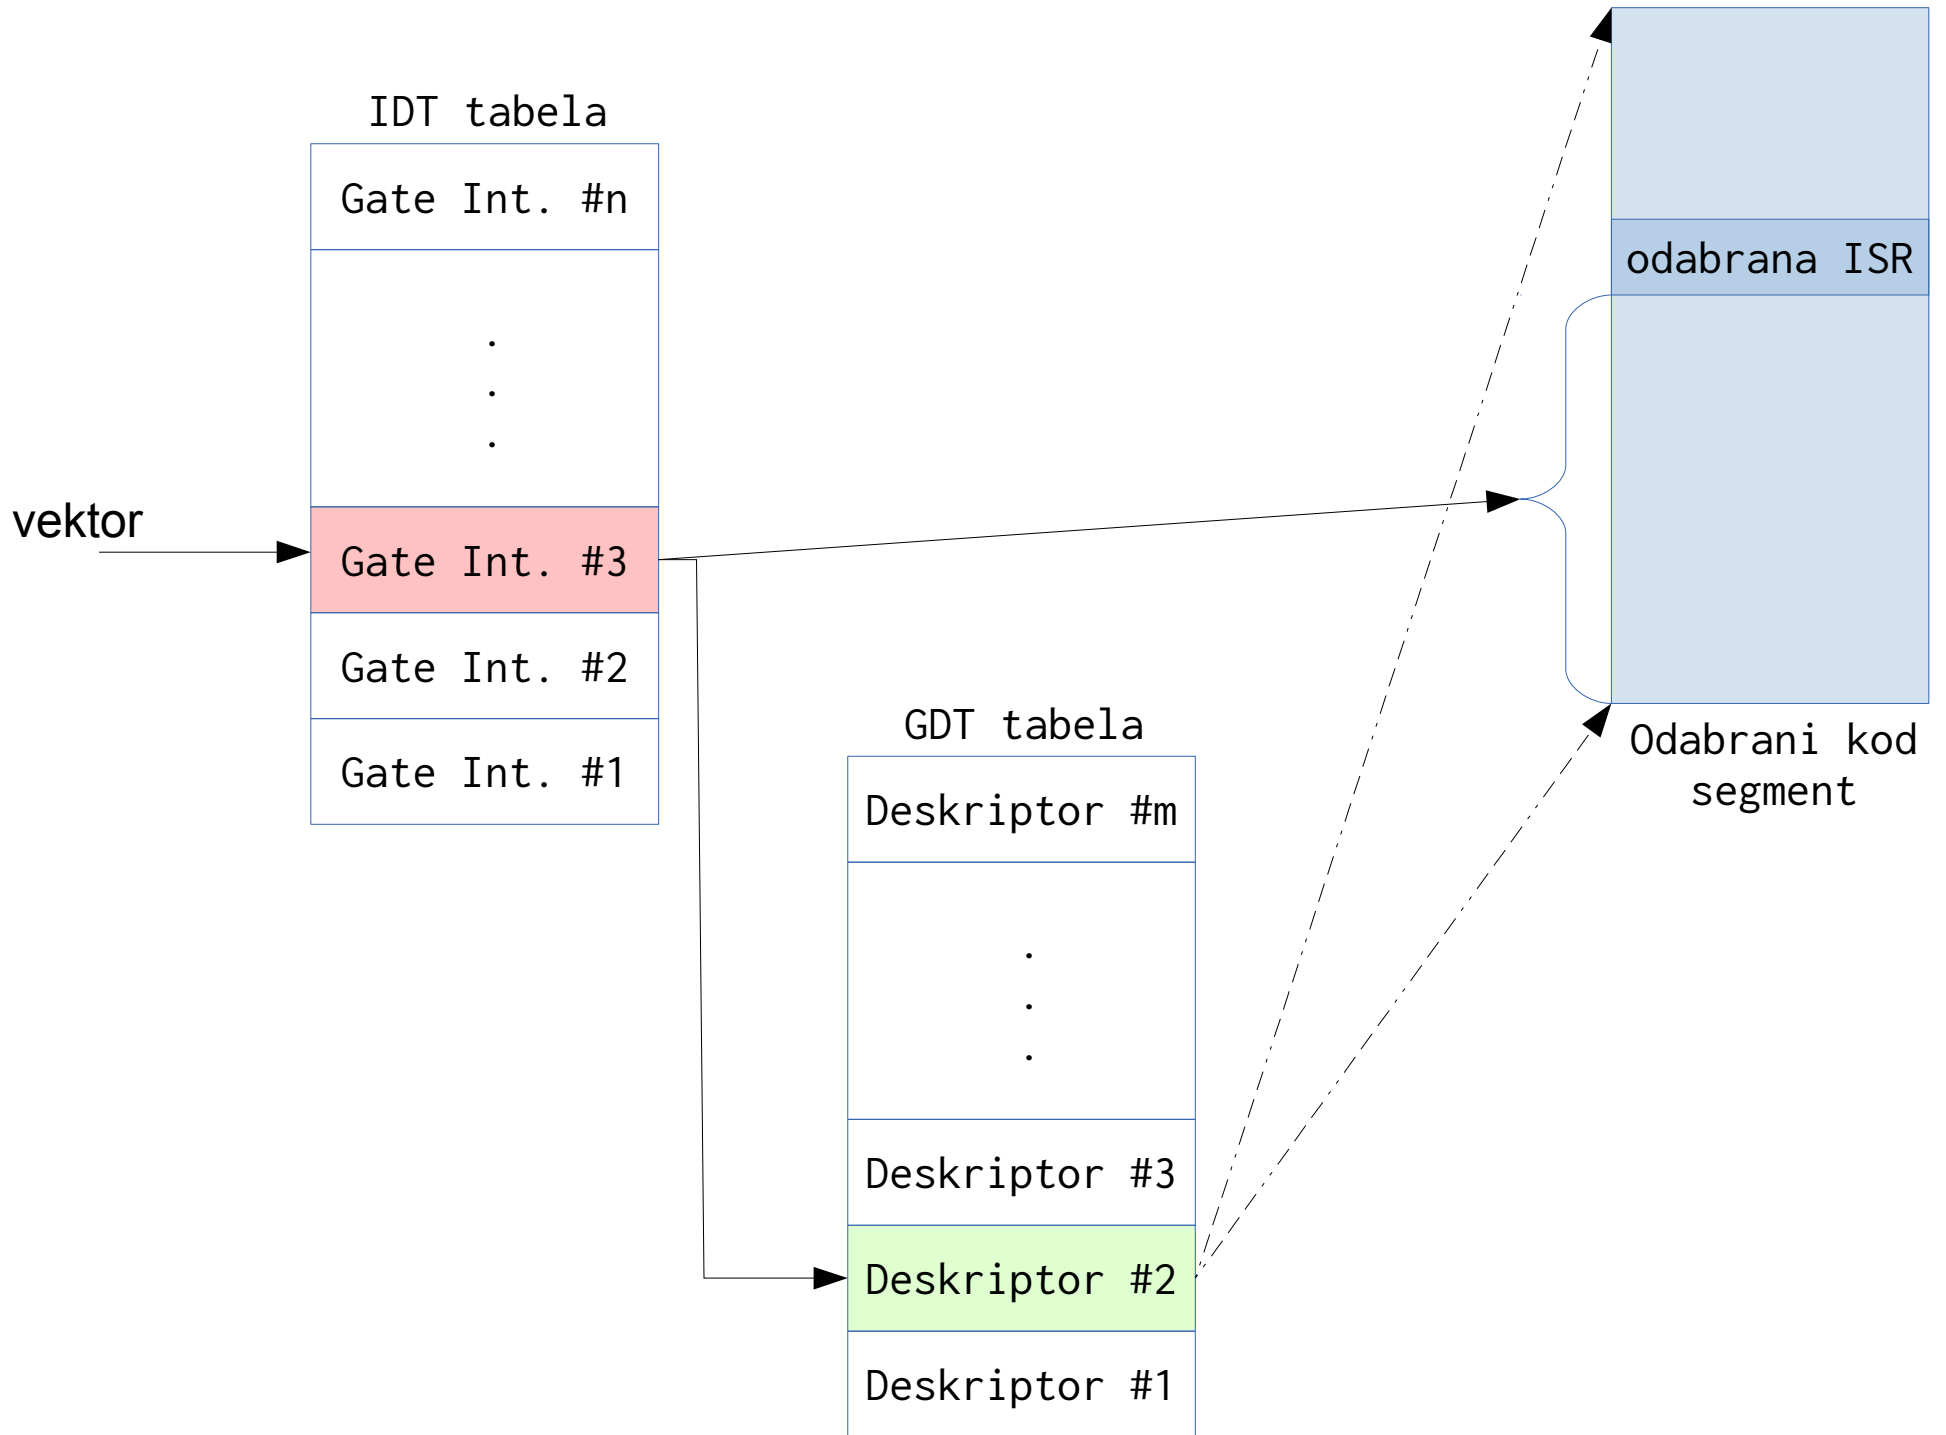
dr.sc. Amer Hasanović

Tretman iznimki i sistemskih poziva

# Vektor prekida

- Svaki prekid, bez obzira na tip, asocira se sa brojem, tzv **vektor**, koji identificira konkretni prekid.
  - vektore 0–32 x86 rezervira za iznimke
  - ostali su konfigurabilni za hardver prekide i systemske pozive
- Kada nastane prekid, na osnovu vrijednosti vektora bira se odgovorajuća rutina koja tretira prekid:
  - vektorom se indeksira IDT (Interrupt Descriptor Table), a iz dobivenog deskriptora (ili gate-a) čita se destinacijski %cs i %eip;
  - adresu IDT-a CPU drži u posebnom registru %idtr, čija se vrijednost mijenja instrukcijom `lidt`
  - operativni sistem u procesu inicijalizacije mora kreirati tabelu koja definira gate-e za sve moguće prekide, te istu učitati u svako aktivno jezgro CPU-a sa `lidt`






# xv6 i IDT


- xv6 inicijalizira IDT u funkciji `tvinit`
- Vrijednost svakog pojedinačnog gate-a unutar IDT postavlja se putem makroa `SETGATE`
- `SETGATE(gate, istrap, sel, off, d)`
  - `gate` – red u IDT tabeli koji se definira
  - `istrap` – primanje hardver prekida tokom tretmana datog prekida (1 → da, 0 → ne)
  - `sel` – asocirani segment deskriptor za dati prekid (%cs)
  - `off` – adresa asoricranog ISR (%eip)
  - `d` – DPL polje unutar gate-a koje definira minimalni nivo privilegija potreban da se prekid asociran sa datim vektorom generira softverski putem `int` instrukcije

- xv6 definira ISR rutine za sve prekide u fajlu `vectors.S`
  - Obratiti pažnju da hardver na steku generira Error code polje za prekide asocirane sa vektorima 8, 10-14, 17, dok za ostale vektore njihovi ISR na stek stavljaju umjetni Error code

Prije prekida CPU u kernel modu:



Prije prekida CPU u user modu:



```

struct gatedesc idt[256];
extern uint vectors[];
void tvinit(void)
{
    int i;

    for(i = 0; i < 256; i++)
        SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);
    SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3, vectors[T_SYSCALL], DPL_USER);

    initlock(&tickslock, "time");
}

```

trap.c

#define T\_SYSCALL 64

```

.data
.globl vectors
vectors:
    .long vector0
    .long vector1
    .long vector2
    .long vector3
    .long vector4
    .long vector5
    .long vector6
    .long vector7
    .long vector8
    ...
    .long vector254
    .long vector255

```

```

.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
.globl vector1
vector1:
    pushl $0
    pushl $1
    jmp alltraps
...
vector8:
    pushl $8
    jmp alltraps

```

```

...
.globl vector254
vector254:
    pushl $0
    pushl $254
    jmp alltraps
.globl vector255
vector255:
    pushl $0
    pushl $255
    jmp alltraps

```



```

.globl alltraps
alltraps:
    # Build trap frame.
    pushl %ds
    pushl %es
    pushl %fs
    pushl %gs
    pushal

    # Set up data and per-cpu segments.
    movw $(SEG_KDATA<<3), %ax
    movw %ax, %ds
    movw %ax, %es
    movw $(SEG_KCPU<<3), %ax
    movw %ax, %fs
    movw %ax, %gs

    # Call trap(tf), where tf=%esp
    pushl %esp
    call trap
    addl $4, %esp

    # Return falls through to trapret...
.globl trapret
trapret:
    popal
    popl %gs
    popl %fs
    popl %es
    popl %ds
    addl $0x8, %esp # trapno and errcode
    iret

```

ss
esp
eflags
cs
eip
err. code
trapno

```
.globl alltraps
```

```
alltraps:
```

```
    # Build trap frame.
```

```
    pushl %ds
```

```
    pushl %es
```

```
    pushl %fs
```

```
    pushl %gs
```

```
    pushal
```

```
    # Set up data and per-cpu segments.
```

```
    movw $(SEG_KDATA<<3), %ax
```

```
    movw %ax, %ds
```

```
    movw %ax, %es
```

```
    movw $(SEG_KCPU<<3), %ax
```

```
    movw %ax, %fs
```

```
    movw %ax, %gs
```

```
    # Call trap(tf), where tf=%esp
```

```
    pushl %esp
```

```
    call trap
```

```
    addl $4, %esp
```

```
    # Return falls through to trapret...
```

```
.globl trapret
```

```
trapret:
```

```
    popal
```

```
    popl %gs
```

```
    popl %fs
```

```
    popl %es
```

```
    popl %ds
```

```
    addl $0x8, %esp # trapno and errcode
```

```
    iret
```

ss
esp
eflags
cs
eip
err. code
trapno
ds
es
fs
gs

```

.globl alltraps
alltraps:
    # Build trap frame.
    pushl %ds
    pushl %es
    pushl %fs
    pushl %gs
    pushal

    # Set up data and per-cpu segments.
    movw $(SEG_KDATA<<3), %ax
    movw %ax, %ds
    movw %ax, %es
    movw $(SEG_KCPU<<3), %ax
    movw %ax, %fs
    movw %ax, %gs

    # Call trap(tf), where tf=%esp
    pushl %esp
    call trap
    addl $4, %esp

    # Return falls through to trapret...
.globl trapret
trapret:
    popal
    popl %gs
    popl %fs
    popl %es
    popl %ds
    addl $0x8, %esp # trapno and errcode
    iret

```

ss
esp
eflags
cs
eip
err. code
trapno
ds
es
fs
gs
eax
ecx
edx
ebx
esp
ebp
esi
edi



```

.globl alltraps
alltraps:
    # Build trap frame.
    pushl %ds
    pushl %es
    pushl %fs
    pushl %gs
    pushal

    # Set up data and per-cpu segments.
    movw $(SEG_KDATA<<3), %ax
    movw %ax, %ds
    movw %ax, %es
    movw $(SEG_KCPU<<3), %ax
    movw %ax, %fs
    movw %ax, %gs

    # Call trap(tf), where tf=%esp
    pushl %esp
    call trap
    addl $4, %esp

    # Return falls through to trapret...
.globl trapret
trapret:
    popal
    popl %gs
    popl %fs
    popl %es
    popl %ds
    addl $0x8, %esp # trapno and errcode
    iret

```

ss
esp
eflags
cs
eip
err. code
trapno
ds
es
fs
gs
eax
ecx
edx
ebx
esp
ebp
esi
edi
esp



```

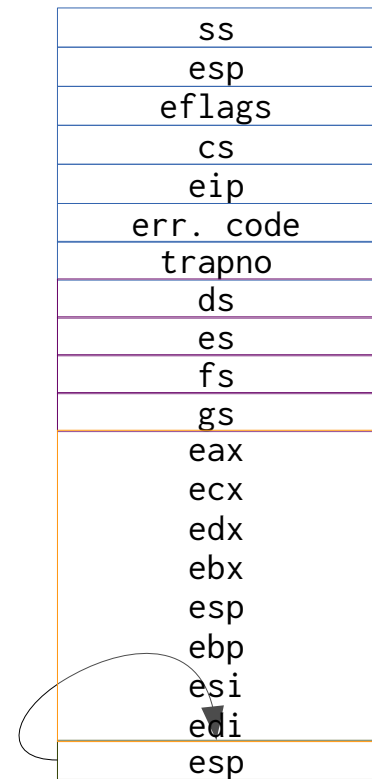
.globl alltraps
alltraps:
    # Build trap frame.
    pushl %ds
    pushl %es
    pushl %fs
    pushl %gs
    pushal

    # Set up data and per-cpu segments.
    movw $(SEG_KDATA<<3), %ax
    movw %ax, %ds
    movw %ax, %es
    movw $(SEG_KCPU<<3), %ax
    movw %ax, %fs
    movw %ax, %gs

    # Call trap(tf), where tf=%esp
    pushl %esp
    call trap
    addl $4, %esp

    # Return falls through to trapret...
.globl trapret
trapret:
    popal
    popl %gs
    popl %fs
    popl %es
    popl %ds
    addl $0x8, %esp # trapno and errcode
    iret

```



```

struct trapframe {

    uint edi;
    uint esi;
    uint ebp;
    uint oesp;      // useless & ignored
    uint ebx;
    uint edx;
    uint ecx;
    uint eax;


    ushort gs;
    ushort padding1;
    ushort fs;
    ushort padding2;
    ushort es;
    ushort padding3;
    ushort ds;
    ushort padding4;
    uint trapno;


    uint err;
    uint eip;
    ushort cs;
    ushort padding5;
    uint eflags;


    uint esp;
    ushort ss;
    ushort padding6;
};

```

ss
esp
eflags
cs
eip
err. code
trapno
ds
es
fs
gs
eax
ecx
edx
ebx
esp
ebp
esi
edi

```
void trap(struct trapframe *tf)
```

```
{  
    if(tf->trapno == T_SYSCALL) {  
        if(proc->killed)  
            exit();  
        proc->tf = tf;  
        syscall();  
        if(proc->killed)  
            exit();  
        return;  
    }
```

} tretman sistemskih poziva

```
  
    switch(tf->trapno) {  
    case T_IRQ0 + IRQ_TIMER:  
        if(cpu->id == 0){  
            acquire(&tickslock);  
            ticks++;  
            wakeup(&ticks);  
            release(&tickslock);  
        }  
        lapiceoi();  
        break;  
    case T_IRQ0 + IRQ_IDE:  
        ideintr();  
        lapiceoi();  
        break;  
    case T_IRQ0 + IRQ_IDE+1:  
        // Bochs generates spurious IDE1 interrupts.  
        break;  
    case T_IRQ0 + IRQ_KBD:  
        kbdintr();  
        lapiceoi();  
        break;  
    case T_IRQ0 + IRQ_COM1:  
        uartintr();  
        lapiceoi();  
        break;  
    case T_IRQ0 + 7:  
    case T_IRQ0 + IRQ_SPURIOUS:  
        cprintf("cpu%d: spurious interrupt at %x:%x\n",  
                cpu->id, tf->cs, tf->eip);  
        lapiceoi();  
        break;  
    }
```

} tretman hardver prekida

```
// nastavak sa prethodnog slajda
```

```
default:
```

```
if(proc == 0 || (tf->cs&3) == 0) {
```

```
    // In kernel, it must be our mistake.
```

```
    cprintf("unexpected trap %d from cpu %d eip %x (cr2=0x%x)\n",
```

```
        tf->trapno, cpu->id, tf->eip, rcr2());
```

```
    panic("trap");
```

```
}
```

```
// In user space, assume process misbehaved.
```

```
cprintf("pid %d %s: trap %d err %d on cpu %d "
```


```
    "eip 0x%x addr 0x%x--kill proc\n",
```

```
    proc->pid, proc->name, tf->trapno, tf->err, cpu->id, tf->eip,
```

```
    rcr2());
```

```
proc->killed = 1;
```

```
}
```



tretman iznimki



# sistemiški pozivi

- xv6 koristi vektor 64 (T\_SYSCALL) za signalizaciju sistemskog poziva.
- ukoliko neka aplikacija želi da pozove neku od kernel funkcija eksportovanih u korisnički prostor kroz mehanizam sistemskih poziva, aplikacija treba da obavi korake:
  - broj sistemskog poziva koji xv6 asocira sa željenom kernel funkcijom postaviti u registar %eax;
  - argumente za kernel funkciju postaviti na stek aplikacije;
  - izvršiti sistemski poziv sa instrukcijom: `int 64`

- Svaki sistemski poziv ima dva dijela:
  - funkciju za pripremu poziva unutar korisničkog prostora:
    - kompajlira se i linka skupa sa aplikacijom;
    - xv6 relevantni fajlovi: `user.h` i `usys.S`
  - funkciju koja predstavlja implementaciju sistemskog poziva
    - ove funkcije kompajliraju se i linkaju sa ostalim kernel funkcijama;
    - svaka funkcija treba da dobije identifikator koji predstavlja broj sistemskog poziva;
    - xv6 relevantni fajlovi: `syscall.c` i `sysproc.c`

# Sigurnost sistemskog poziva

- Argumente za sistemski poziv aplikacija ostavlja na svom steku.
  - kada se aktivira kernel nakon prekida kod sistemskog poziva, on ima pristup i kernel i user dijelu AP-a procesa;
  - prije upotrebe argumenata, kernel mora provjeriti validnost istih jer eventualna iznimka (npr page fault) u kernel prostoru može izazvati zaustavljanje cijelog OS-a
    - xv6 za provjeru argumenata koristi funkcije: argint, argptr, argstr itd..

```
int argint(int n, int *ip)
{
    return fetchint(proc->tf->esp + 4 + 4*n, ip);
}
```

```
int fetchint(uint addr, int *ip)
{
    if(addr >= proc->sz || addr+4 > proc->sz)
        return -1;
    *ip = *(int*)(addr);
    return 0;
}
```

## user.h

```
// system calls
int fork(void);
int exit(void) __attribute__((noreturn));
int wait(void);
int pipe(int*);
int write(int, void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
int exec(char*, char**);
int open(char*, int);
int mknod(char*, short, short);
int unlink(char*);
int fstat(int fd, struct stat*);
int link(char*, char*);
int mkdir(char*);
int chdir(char*);
int dup(int);
int getpid(void);
char* sbrk(int);
int sleep(int);
int uptime(void);
```

## usys.S

```
.globl fork; fork: movl $1, %eax; int $64; ret
.globl exit; exit: movl $2, %eax; int $64; ret
.globl wait; wait: movl $3, %eax; int $64; ret
.globl pipe; pipe: movl $4, %eax; int $64; ret
.globl read; read: movl $5, %eax; int $64; ret
.globl write; write: movl $16, %eax; int $64; ret
.globl close; close: movl $21, %eax; int $64; ret
.globl kill; kill: movl $6, %eax; int $64; ret
.globl exec; exec: movl $7, %eax; int $64; ret
.globl open; open: movl $15, %eax; int $64; ret
.globl mknod; mknod: movl $17, %eax; int $64; ret
.globl unlink; unlink: movl $18, %eax; int $64; ret
.globl fstat; fstat: movl $8, %eax; int $64; ret
.globl link; link: movl $19, %eax; int $64; ret
.globl mkdir; mkdir: movl $20, %eax; int $64; ret
.globl chdir; chdir: movl $9, %eax; int $64; ret
.globl dup; dup: movl $10, %eax; int $64; ret
.globl getpid; getpid: movl $11, %eax; int $64; ret
.globl sbrk; sbrk: movl $12, %eax; int $64; ret
.globl sleep; sleep: movl $13, %eax; int $64; ret
.globl uptime; uptime: movl $14, %eax; int $64; ret
```

syscall.c

```
void syscall(void)
{
    int num;
    num = proc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        proc->tf->eax = syscalls[num]();
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            proc->pid, proc->name, num);
        proc->tf->eax = -1;
    }
}
```

#define SYS\_getpid 11

```
static int (*syscalls[])(void) = {
    [SYS_fork]      sys_fork,
    [SYS_exit]      sys_exit,
    [SYS_wait]      sys_wait,
    [SYS_pipe]      sys_pipe,
    [SYS_read]      sys_read,
    [SYS_kill]      sys_kill,
    [SYS_exec]      sys_exec,
    [SYS_fstat]     sys_fstat,
    [SYS_chdir]     sys_chdir,
    [SYS_dup]       sys_dup,
    [SYS_getpid]    sys_getpid,
    [SYS_sbrk]      sys_sbrk,
    [SYS_sleep]     sys_sleep,
    [SYS_uptime]    sys_uptime,
    [SYS_open]      sys_open,
    [SYS_write]     sys_write,
    [SYS_mknod]     sys_mknod,
    [SYS_unlink]    sys_unlink,
    [SYS_link]      sys_link,
    [SYS_mkdir]     sys_mkdir,
    [SYS_close]     sys_close,
};
```

sysproc.c

```
int sys_getpid(void)
{
    return proc->pid;
}
```